

# Exam 1 - Introduction to Data Science - Fall 2021

## Guidelines

- Answer in the Markdown or code cells given below the questions.
- Feel free to use the lecture notes and other resources but work on your own!
- Comment code where appropriate for clarity.
- Use only packages that are loaded with the code in a given problem.

## Problem 1

### 1.1

We load *lowbt* for you.

```
In [1]: # Scikit-learn function used to load dataset from OpenML.  
from sklearn.datasets import fetch_openml  
  
# Load 'lowbwt' bunch from OpenML.  
lowbwt_bunch = fetch_openml('lowbwt', version=1, as_frame=True)  
# Extract dataset from bunch.  
lowbwt = lowbwt_bunch.data  
display(lowbwt)
```

	LOW	AGE	LWT	RACE	SMOKE	PTL	HT	UI	FTV
0	0	19.0	182.0	2	0	0	0	1	0
1	0	33.0	155.0	3	0	0	0	0	3
2	0	20.0	105.0	1	1	0	0	0	1
3	0	21.0	108.0	1	1	0	0	1	2
4	0	18.0	107.0	1	1	0	0	1	0
...	...	...	...	...	...	...	...	...	...
184	1	28.0	95.0	1	1	0	0	0	2
185	1	14.0	100.0	3	0	0	0	0	2
186	1	23.0	94.0	3	1	0	0	0	0
187	1	17.0	142.0	2	0	0	1	0	0
188	1	21.0	130.0	1	1	0	1	0	3


189 rows × 9 columns

Plot a **density histogram** for the ages of mothers with no uterine irritability.

In [17]:

```
import pandas as pd
import matplotlib.pyplot as plt

# YOUR CODE HERE.
indices = (lowbwt["UI"] == "0")
mothers_no_ui = lowbwt["AGE"][indices]
mothers_no_ui.plot.density(title='Density histogram for the ages of mothers w
plt.show()
```



## 1.2

A student runs the following code to test the behavior of the `scipy.stats` i.i.d. random variate generation for a standard normal  $\mathcal{N}(0, 1)$  distribution. He wants to check whether the variance of a randomly generated i.i.d. sample of size ten will be close to the theoretical variance of 1 on average. To accomplish this, he

1. Simulates 10,000 random i.i.d. samples of size ten.
2. Takes the variances of each of the 10,000 samples.
3. Averages all of these variances, and prints the average.

When he does this, he gets an average of variance that is close to 0.90 instead of 1.00. He wrongly concludes that `scipy.stats` is broken. Change one line of code so that the average variance will be as expected.

In [7]:

```
import scipy.stats as stats
import numpy as np

# CORRECT ONE LINE IN THE FOLLOWING CODE

# Generate 10,000 random samples of size ten.
samples = stats.norm.rvs(size=[10,10000], random_state=553)

# Take variance of each sample for unbiased estimation
sample_variances = np.var(samples, axis=1, ddof=1)

# Take average of variances.
average_variance = np.mean(sample_variances)
# Print.
print(average_variance)
```

0.9974394297910786

## Problem 2

## 2.1

The empirical cumulative distribution function (ECDF)  $\hat{F}_n(t)$  of the samples  $\{x_1, x_2, \dots, x_n\}$  is defined by

$$\hat{F}_n(t) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}_{x_i \leq t} = \frac{\#\{i \text{ such that } x_i \leq t\}}{n},$$

Here,  $\mathbf{1}_{x_i \leq t}$  is equal to 1 if  $x_i \leq t$  and equal to 0 otherwise.

We load the samples to the numpy array `samples` below.

```
In [1]: import numpy as np
samples = np.array([
    -0.986884764999009, -0.851970773256441, -0.867796672206323, 1.16345615827
    -1.09025113310314, -1.00814879621656, 0.806680833662076, 1.06986054952863
    -1.12102128751188, 1.04386395472623, -0.881116187067203, -0.9419772708833
])
```


Manually implement an empirical CDF function of the provided samples  $\{x_1, \dots, x_n\}$ , then run the cell to plot the ECDF of the samples.

```
In [2]: import numpy as np
import matplotlib.pyplot as plt

def my_ecdf(vec, dat):
    # YOUR CODE HERE.
    cdf_vals = []
    for i in range(len(vec)):
        count = sum(dat < vec[i])
        cdf_vals.append(count / len(dat))

    return cdf_vals

horiz_axis = np.linspace(min(samples), max(samples), 100)
plt.plot(horiz_axis, my_ecdf(horiz_axis, samples))
plt.show()
```



## 2.2

The following code is an attempt to implement triangular kernel density estimation from the sample  $\{x_1, \dots, x_n\}$  provided by `samples`, where the triangular kernel  $K(x)$  is defined by

$$K(x) = \begin{cases} 1 - |x| & \text{if } |x| \leq 1 \\ 0 & \text{otherwise} \end{cases}$$

In the first cell, we load a sample to the numpy array `samples`.

```
In [3]: import numpy as np

# Load sample.
samples = np.array([
    -0.986884764999009, -0.851970773256441, -0.867796672206323, 1.16345615827
    -1.09025113310314, -1.00814879621656, 0.806680833662076, 1.06986054952863
    -1.12102128751188, 1.04386395472623, -0.881116187067203, -0.9419772708833
])
```

In the following cell, there is an error either in the function definition of `kernel` or of `density` which causes the integration of the KDE to be not close to 1. Correct the error.

In [4]:

```

import numpy as np
import matplotlib.pyplot as plt

# CORRECT THE FOLLOWING CODE


# Triangular kernel function.
def kernel(x, dat, bw):
    t = 1 - abs((x - dat) / bw)
    return(t)

# Kernel density estimator function.
def density(x, sample, bw):
    f = np.zeros(len(x))
    for dat in sample:
        f = f + kernel(x, dat, bw)
    return(f / (len(sample)*bw))

# Plot kernel density estimate for sample.
x = np.linspace(-3,3, num = 500)
p = density(x, samples, 0.25) / np.trapz(density(x, samples, 0.25), x)
plt.plot(x, p)
plt.show()

# Check integral of kernel density estimate.
np.trapz(p, x)

```



Out[4]: 1.0

## Problem 3

## 3.1

Use the [SciPy stats library](#) to generate a random sample of size 10,000 from an [exponential distribution](#) with `scale=1/15`, which means that the parameter  $\lambda = 15$ . Use a `random_state` of 553. Plot a density histogram of the result with 100 bins, and overlay the density of the distribution.

In [17]:


```
import seaborn as sns
import numpy as np
import matplotlib.pyplot as plt
%matplotlib inline

sns.set(rc={'figure.figsize':(12.5, 9.5)})
sns.set_context('talk')

np.random.seed(553) # seed random number generator with fixed value so we always
exponential_distribution_values = list(np.random.exponential(scale=1/15, size=10000))

sns.distplot(exponential_distribution_values, bins = 100, kde=False, color='darkblue')
```

Out[17]: <AxesSubplot:>



## 3.2

The following code is an attempt to implement rejection sampling from a [Student's  \$t\$ -distribution](#) with  $df=5$ . More precisely, the target distribution is supposed to be a Student's  $t$ -distribution with  $df=5$ , and the candidate distribution is supposed to be a standard normal distribution (a  $\mathcal{N}(0, 1)$  distribution).

As you can see, the output is off. Change one line of code to produce the correct output.



In [20]:

```
import scipy.stats as stats
import matplotlib.pyplot as plt


# CORRECT ONE LINE IN THE FOLLOWING CODE

trials = 100000
df = 5
U = stats.uniform.rvs(loc=0, scale=1, size=trials, random_state=436)
Y = stats.norm.rvs(size=trials, random_state=636)

x = np.linspace(-4, 4, 1000)
M = max(stats.t.pdf(x, df=df) / stats.norm.pdf(x))
f = stats.t.pdf(Y, df=df)
g = stats.norm.pdf(Y)
ratio = f/g
samples = Y[U <= ratio/M]

print('Number of samples:\t', len(samples))
plt.hist(samples, density=True, bins=25)
plt.plot(x, stats.norm.pdf(x), label='norm')
plt.plot(x, stats.t.pdf(x, df=df), label='t')
plt.legend()
plt.show()
```

Number of samples: 2494



## Problem 4

### 4.1

In the following cell, we load the *boston* dataset, store its feature names in the variable `labels`, and divide it into a training and test set.

```
In [8]: from sklearn.datasets import load_boston
from sklearn.model_selection import train_test_split

boston_bunch = load_boston()
X = boston_bunch.data
y = boston_bunch.target
labels = boston_bunch.feature_names

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, ran
```

1. Perform Lasso regression on all predictor variables in the training set to predict housing price. Choose the largest alpha in the interval  $[0, 1]$ , to the nearest tenth, at which the test R-squared of the resulting predictions is above 0.65.
2. Print the alpha that you choose along with the resulting R-squared.
3. Print a list of the names of the variables positively associated with housing price and a list of the names of the variables negatively associated with housing price.

In [ ]:

```
import numpy as np
from sklearn.linear_model import Lasso

alpha =                # the regularization parameter alpha

#YOUR CODE HERE.

R_squared =            # the R-squared

Positive_features =     # the names of the variables positively associate
Negative_features =     # the names of the variables negatively associate
print('{:<15} {}'.format('Alpha:', alpha))
print('{:<15} {}'.format('R-squared:', R_squared))
print('{:<15} {}'.format('Positive:', Positive_features))
print('{:<15} {}'.format('Negative:', Negative_features))
```

## 4.2

A student decides to select a predictive model for housing price on the basis of MSE. He runs the following code to fit three different models.

In [9]:

```

import numpy as np
from sklearn.linear_model import LinearRegression
from sklearn.metrics import mean_squared_error
import sklearn.preprocessing as prepro

# MODEL 1: using all predictor variables
X_train_1, X_test_1 = X_train, X_test
reg_1 = LinearRegression().fit(X_train_1, y_train)

# MODEL 2: using only AGE, NOX, DIS, and RAD as predictor variables
ind = np.where([a in ['ZN', 'RM', 'AGE', 'RAD', 'B'] for a in boston_bunch.feature_names])[0]
X_train_2, X_test_2 = X_train[:,ind], X_test[:,ind]
reg_2 = LinearRegression().fit(X_train_2, y_train)

# MODEL 3: using all polynomial combinations of degree <=5 of the original th
poly = prepro.PolynomialFeatures(5)
X_train_3, X_test_3 = poly.fit_transform(X_train), poly.fit_transform(X_test)
reg_3 = LinearRegression().fit(X_train_3, y_train)

```

Based on the following code and output, he selects MODEL 3 for having the best MSE. The student made a mistake on one line: correct that line.

In [10]:

```

# CORRECT THE FOLLOWING CODE IN ONE LINE

# Function for comparing MSE of models on data.
def model_comparison(model_list, features_list, target):
    k = 0
    for model in model_list:
        # Calculate MSE.
        mse = mean_squared_error(target, model.predict(features_list[k]))
        # Print model name.
        print('MODEL', k+1)
        # Print MSE.
        print('{:<15} {}'.format('MSE:', mse))
        k = k+1
    return

model_list = [reg_1, reg_2, reg_3]
features_list, target = [X_train_1, X_train_2, X_train_3], y_train
model_comparison(model_list, features_list, target)

```

```

MODEL 1
MSE:          18.637149771206403
MODEL 2
MSE:          32.16123650012168
MODEL 3
MSE:          3.0771694016607596e-19

```

## Problem 5

## 5.1

We load the [following dataset](#) on 310 orthopedic patients, including 6 predictors concerning the orientation of bones in the lower back, as well as a target class variable for normal ( 2 ) or abnormal ( 1 ) spinal health. We scale the predictors.

```
In [46]: import pandas as pd
import sklearn
from sklearn.preprocessing import StandardScaler

vertebra = pd.read_csv('https://www.openml.org/data/get_csv/1593720/phpzeLjnh')
predictors = vertebra.iloc[:,0:6]
target = vertebra.iloc[:,6]

print(vertebra)

scaler = StandardScaler()
predictors_scaled = scaler.fit_transform(predictors)

print(predictors_scaled)
```

	V1	V2	V3	V4	V5	V6	Class
0	63.027818	22.552586	39.609117	40.475232	98.672917	-0.254400	1
1	39.056951	10.060991	25.015378	28.995960	114.405425	4.564259	1
2	68.832021	22.218482	50.092194	46.613539	105.985136	-3.530317	1
3	69.297008	24.652878	44.311238	44.644130	101.868495	11.211523	1
4	49.712859	9.652075	28.317406	40.060784	108.168725	7.918501	1
..	...	...	...	...	...	...	...
305	47.903565	13.616688	36.000000	34.286877	117.449062	-4.245395	2
306	53.936748	20.721496	29.220534	33.215251	114.365845	-0.421010	2
307	61.446597	22.694968	46.170347	38.751628	125.670725	-2.707880	2
308	45.252792	8.693157	41.583126	36.559635	118.545842	0.214750	2
309	33.841641	5.073991	36.641233	28.767649	123.945244	-0.199249	2

```
[310 rows x 7 columns]
[[ 0.14708639  0.50136874 -0.6651769  -0.18495027 -1.44764679 -0.70805942]
 [-1.24586434 -0.74876903 -1.45300076 -1.04152066 -0.26438491 -0.57955636]
 [ 0.48436951  0.46793217 -0.09926173  0.2730833  -0.89768553 -0.79542094]
 ...
 [ 0.05520139  0.51561809 -0.3109775  -0.31356367  0.58289258 -0.77348835]
 [-0.88582308 -0.88565955 -0.55861261 -0.47712773  0.0470211  -0.69554822]
 [-1.5489268  -1.24785958 -0.82539422 -1.05855697  0.45311697 -0.70658867]]
```

Compute and store the PCA transformation of the scaled predictors using all principal components.

```
In [58]: from sklearn.decomposition import PCA

# YOUR CODE HERE
pca = PCA(n_components=6)
pca.fit(predictors_scaled)
Proj_predictors = pca.fit_transform(predictors_scaled) # the PCA transform
```

## 5.2

How many PC's should be retained to account for 80% of the variance in the scaled predictors? Justify your answer.

```
In [57]: import numpy as np
from sklearn.decomposition import PCA

# YOUR CODE HERE
for n_components in range(1, 7):
    pca = PCA(n_components=n_components)
    pca.fit(predictors_scaled)
    print(sum(pca.explained_variance_ratio_))
```

```
0.540963547073773
0.7400606441379418
0.8669090632480623
0.9456637267445556
0.9999999999999998
0.9999999999999999
```

**Answer:**

3